



# ECE 411 - MP4

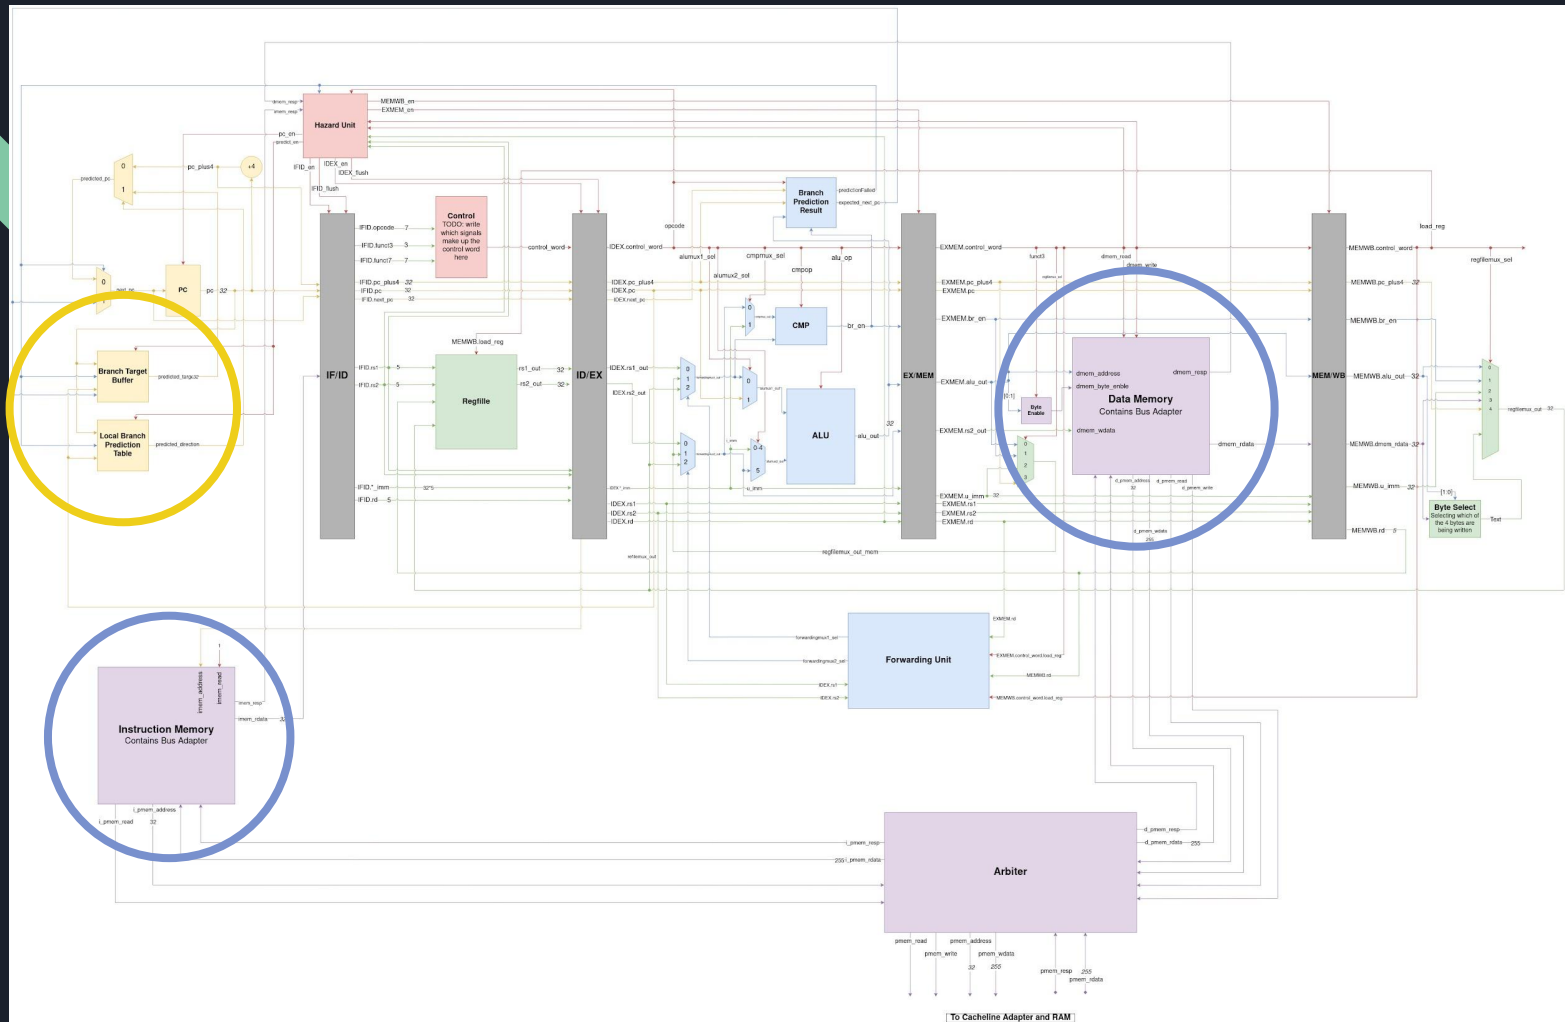
## Final Presentation

**Team: LMP**  
Lakshya, Mridul, Peter



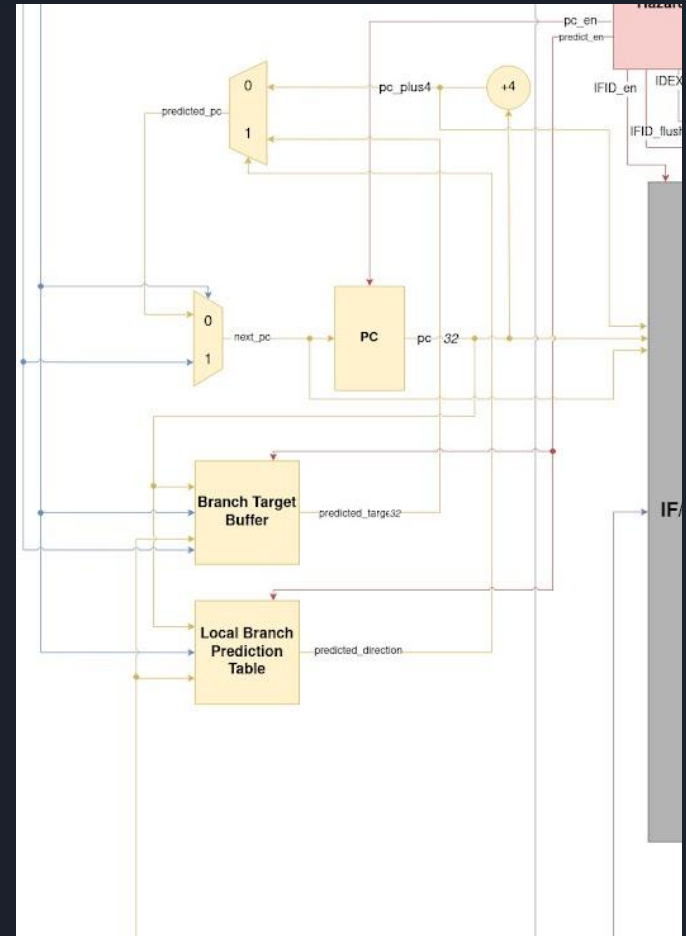
# Feature Overview

- 5-stage Pipelined CPU with Hazard Detection and Forwarding
- Parameterized Local Branch Prediction Table and Branch Target Buffer
- 2-stage Pipelined & Parameterized Split Cache up to 64 Sets and 4 Ways
- RISC-V M-Extension with Wallace Tree Multiplier and Non-Storing Divider
- L2 Cache System (Not included in the Competition)
- Global Branch History Register (Not included in Competition)



# Feature #1: Parameterized Local Branch Prediction Table and Branch Target Buffer

- Single-cycle
  - Tried to use BRAM but this affected functionality due to 2-cycle latency
- Parametrized for Number of Sets
  - Tested for 128-512 sets
- **Added: Checking for BR/Jumps**
- Hashed using lower order bits of PC
  - Tried using other bits
  - 7-bit Global History Register
    - Goal: increased contextual information for BR/Jump
    - Too much 'randomness' in indexing





# Feature #1 Results

Metric	Basic Prediction		Checking Instruction Opcode	
	comp1.s	comp3.s	comp1.s	comp3.s
<b>Clock Cycles</b>	55,885	72,665	49,537	68,835
BR Accuracy	84.0 %	56.6 %	89.3 %	67.5 %
J Accuracy	99.9 %	~100.0 %	99.9 %	~100.0 %
<b>Total Accuracy</b>	<b>87.8 %</b>	<b>75.1 %</b>	<b>91.9 %</b>	<b>81.8 %</b>

About 4.65% and 8.95% improvement respectively



## Feature #2: 2-stage Pipelined & Parameterized Split Cache up to 64 Sets and 4 Ways

- Pipelining
  - Modified the state machine to be able to receive new memory requests while responding to the previous requests
  - Stores previous request's address and data into registers
  - No need to stall for consecutive hits
- Parameterized Cache
  - Parameterized the number of index bits, cacheline size, and number of ways
  - Used the Pseudo LRU design from lecture
  - Tested with up to 64 sets and 4 ways
- BRAM
  - Used BRAM for reduced power and resource utilization

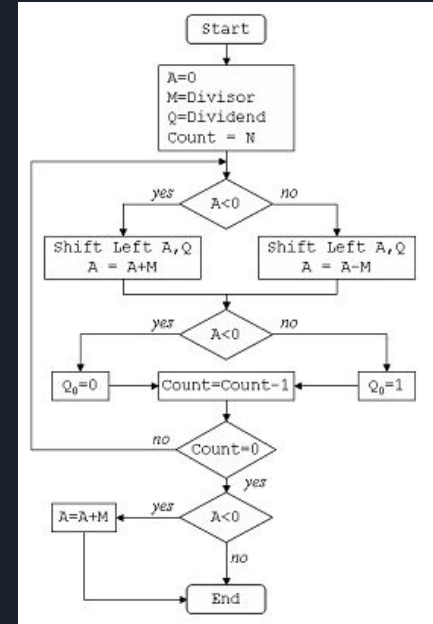
## Feature #2 Results

Metric	8 Sets, 2 Ways			16 Sets, 4 Ways			32 Sets, 4 Ways		
	comp1.s	comp2_i.s	comp3.s	comp1.s	comp2_i.s	comp3.s	comp1.s	comp2_i.s	comp3.s
<b>Clock Cycles</b>	50,377	385,383	341,533	50,057	237,319	72,665	50,057	119,570	72,665
I-Cache Hit Rates	99.9 %	91.0%	82.8%	99.9%	95.6%	99.9%	99.9%	99.9%	99.9%
D-Cache Hit Rates	99.7 %	98.9%	89.6%	99.7 %	99.8%	95.2%	99.7%	99.5%	95.2%

Up to 79% improvement in execution time

# Feature #3: RISC-V M-Extension with Wallace Tree Multiplier and Non-Storing Divider

- Wallace Tree Multiplier
  - Implemented a 2-cycle 32 bit unsigned Wallace Tree Multiplier
  - Used 2-cycle because the delay is too long for single-cycle
- Non-Storing Divider
  - Implemented an 8-cycle unsigned Non-Storing Divider
  - The division originally takes 32 cycles, but we fit in 4 division cycles within 1 clock cycle



Divider Logic





## Feature #3 Results

Metrics	comp2_i.s	comp2_m.s
Clock Cycles	119,570	25,924
Execution Time	1,195,755 ns	259,295 ns
# of Instructions	116,080	11,082
Cycles per Instruction	1.03	2.34

About 78% improvement in execution time

# Overall Performance

- 32 sets, 4-way I- and D-cache
- BPT and BTB 128 sets with opcode checking in Fetch stage
- M-Extension

	comp1.s	comp2_m.s	comp3.s
<b>FMax (MHz)</b>	91.11		
<b>Scaled Time (ns)</b>	543,766	287,098	755,576
<b>Power (mW)</b>	607.17	608.16	641.41*
<b>Score</b>	9.7625E-11	1.4391E-11	2.7668E-10
<b>Geometric Mean</b>	<b>7.2981E-11</b>		

## Utilization:

Logic - 20,366 (56%)

Registers - 8181 (23%)

Total Pins - 165 (57%)

M9K Blocks - 68 (21%)

\*This is an upper estimate from Quartus since the VCD file could not be generated due to space limitations.



# Challenges & Optimizations

- Challenges Faced
  - Multiplier causing too much delay
  - Larger BTB led too high resource utilization
  - Weren't able to pipeline BTB to be able to use BRAM
    - BTB was cause for high power and resource utilization
- Tricks to improve performance
  - Removing Priority Encoding (If-Else constructs)
  - Quartus based optimizations
    - Fitter effort → set for Performance
    - Register Retiming
    - Physical Synthesis



# Future Developments

- More sophisticated Branch Prediction
  - Default taken predictions for BR/JAL
  - Tournament Predictor with global history
  - Associative BPT and BTB
    - Remove some aliasing
- 8-Way associative cache
  - $2 \rightarrow 4$  way gave a large improvement